

Event Store Selection for the Janus Backend

Purpose of an Event Store

The Janus architecture uses a slot-based model where each component fills a well-defined role in the stack. A **document store** captures unstructured data and a **relational database** holds normalized state. An **event store** provides a *chronological record* of every significant change in the system. In an event-sourced design, events are the source of truth and allow the system to reconstruct state by replaying events. Event sourcing improves auditability, avoids lost history and simplifies distributed workflows [393156474872709†L42-L56]. Each event is appended to an *append-only log* and is never mutated; consumers build projections or materialized views from the event stream [393156474872709†L84-L104]. The Janus platform needs an event store to:

- * Record all agent actions, system events and user commands to enable *audit trails* and reproducibility.
- * Support *append-only writes* with immutable events to maintain the full history [450349005886131†L120-L124].
- * Provide *efficient retrieval* of events for a particular entity or stream (e.g., a user, project or agent run) [640399355648082†L148-L166].
- * Offer *concurrency control* so that multiple writers cannot overwrite the same stream state [640399355648082†L168-L171].
- * Run *self-hosted* in offline/edge environments, with high availability and security (TLS, authentication, encryption at rest).

Requirements and Evaluation Criteria

1. **Durability & Ordering** – Each event must be persisted in order and remain immutable. The store should support replaying events to reconstruct state.
2. **Write Throughput & Read Performance** – Even though Janus is designed for a single developer initially, the event store must scale to handle many events per second when multiple agents run concurrently. Efficient indexing per stream is important for fast reads [640399355648082†L148-L166].
3. **Concurrency Control** – Optimistic concurrency (expected version) or transactional isolation is needed to avoid race conditions when writing events [640399355648082†L168-L171].
4. **High Availability & Resilience** – Clustered deployment and replication are essential for fault tolerance [450349005886131†L137-L144].
5. **Integration & Ecosystem** – The solution should integrate well with Python and existing stack components (PostgreSQL, RabbitMQ, Redis), and ideally run in containers orchestrated by Kubernetes.
6. **Security & Access Control** – TLS/SSL encryption, authentication, authorization and encryption at rest. A smaller attack surface and support for offline environments are desirable.
7. **Operational Complexity** – Ease of setup and management, especially for a solo developer. Heavy operational burden is undesirable.

Candidate Technologies

EventStoreDB

Description: EventStoreDB is a specialized database built specifically for event sourcing. It stores data as *streams of immutable events* and indexes them by stream ID. Developers can reconstruct an entity's state by replaying its events [450349005886131†L103-L124]. The database supports real-time streaming and clusters for high availability [450349005886131†L129-L144].

Key capabilities

Capability	Notes
------------	-------

```
|---|---|
| **Immutable events & append-only log** | Once appended, events cannot be
modified. This immutability guarantees data integrity and provides a reliable
audit trail [450349005886131†L103-L124] . |
| **Per-stream indexing & high write throughput** | EventStoreDB can handle
~15 000 writes/sec and efficiently fetch events for a stream thanks to indexes
[640399355648082†L148-L166] . |
| **Optimistic concurrency control** | Writers specify the expected version when
appending events; mismatched versions cause the write to fail, preventing lost
updates [640399355648082†L168-L171] . |
| **Real-time subscriptions & streaming** | Clients can use live subscriptions,
catch-up subscriptions or persistent subscriptions to receive events in real
time [640399355648082†L173-L176] . |
| **High availability & clustering** | Supports clustered deployments with
replication for fault tolerance [450349005886131†L137-L144] . |
| **Security** | Provides built-in TLS encryption and authentication
(user/password or certificates) and supports fine-grained access control. |
```

****Pros****

- * Purpose-built for event sourcing with immutability and per-stream indexes.
- * Scales to high write throughput and supports clustering for high availability.
- * Native features like subscriptions and optimistic concurrency simplify application logic.
- * Integration libraries exist for many languages including Python.

****Cons****

- * A standalone database to deploy and manage; adds operational complexity compared to reusing existing Postgres.
- * Though open source, advanced features are available under commercial license; community support is smaller than mainstream databases.
- * Resource hungry relative to small single-node environments.

PostgreSQL as an Event Store

****Description:**** Rather than adopting a specialized database, an event store can be implemented in an existing relational database. Each event is stored in a table with columns for `id`, `stream_id`, `version` and `data` (JSONB) [311468078120902†L462-L485] . Retrieving events for a stream uses a simple SQL query ordered by version [311468078120902†L495-L501] , and inserting a new event increments the version atomically within a transaction [311468078120902†L508-L520] .

****Key capabilities****

```
| Capability | Notes |
|---|---|
| **ACID transactions & durability** | PostgreSQL provides ACID compliance
[450349005886131†L178-L184] . Events can be inserted within a serializable or
repeatable-read transaction to ensure correctness [311468078120902†L518-L539] .
|
| **Simple schema** | A single table with `stream_id`, `version` and JSONB
payload captures the event stream [311468078120902†L462-L485] . Additional
indexes can be added on `stream_id` and `version` for efficient lookups. |
| **Flexible JSON storage** | Event data can be stored as JSONB, allowing events
to contain arbitrary fields [311468078120902†L462-L485] . |
| **Reuses existing infrastructure** | No new database to deploy; uses the
existing PostgreSQL instance already locked in the Janus stack. |
```

****Pros****

- * Reuses the locked-in PostgreSQL, minimizing operational overhead and cost.
- * ACID transactions guarantee consistency and durability; concurrent writes can be serialized.
- * Simplifies backup and security: Postgres supports TLS, row-level permissions and encryption at rest via pgcrypto or full-disk encryption.
- * Good fit for moderate throughput systems where the volume of events is not extremely high; all state is local and offline-friendly.

****Cons****

- * Performance may become a bottleneck at very high write rates; indexes on a large events table could impact query times.
- * Application code must handle optimistic concurrency by reading the current max version and inserting the next version within a transaction
【311468078120902†L508-L520】.
- * Lacks built-in subscription mechanism; event distribution must be implemented using the message broker (RabbitMQ) or Postgres LISTEN/NOTIFY.

Apache Kafka

****Description:**** Kafka is a distributed messaging platform that stores data in append-only logs (topics) and is optimized for throughput. It is often used for event streaming and integrates with frameworks like Kafka Streams.

****Key capabilities****

Capability	Notes
------------	-------

---	---
-----	-----

High throughput & scalability	Kafka can ingest hundreds of thousands of messages per second and supports horizontal scaling through partitions
--	--

【640399355648082†L184-L200】	.
-----------------------------	---

Durable storage & replication	Topics are replicated across brokers to ensure durability and availability.
--	---

Rich ecosystem	Mature tooling (Kafka Streams, Schema Registry, ksqlDB) and broad community support.
---------------------------	--

****Limitations for event sourcing****

- * ****Limited per-stream indexing**** – Kafka partitions messages by key; retrieving all events for a specific stream requires scanning a partition, which becomes slow as the number of events grows 【640399355648082†L190-L204】.

- * ****No optimistic concurrency control**** – Kafka cannot reject writes based on version; concurrency must be handled externally 【640399355648082†L218-L223】.

- * ****Operational complexity**** – Requires running a Kafka cluster (and often ZooKeeper) and managing partitions, retention policies and consumer groups.

- * ****Resource heavy**** – Suitable for high-scale streaming systems but overkill for moderate workloads.

NATS JetStream

****Description:**** JetStream extends the NATS messaging system with durable storage, message replay and streaming semantics. It offers low-latency pub/sub with persistent streams and is designed for microservices and real-time analytics 【129778631843866†L33-L38】.

****Key capabilities****

Capability	Notes
------------	-------

---	---
-----	-----

Durable message storage & replay	JetStream stores messages in file- or
---	---------------------------------------

memory-based streams and allows consumers to replay events from any point in time [129778631843866†L46-L55] . |
****Delivery semantics****	Supports at-least-once and exactly-once delivery, durable subscriptions and ordered messages [129778631843866†L49-L55] .
****High performance & low latency****	JetStream is optimized for low latency (<20 ms) and high throughput for both persistent and non-persistent messages [129778631843866†L42-L45] .
****High availability & scalability****	Clustering and manual sharding enable horizontal scaling; streams can be sharded across nodes [129778631843866†L52-L107] .
****Security****	Provides TLS for secure communication and fine-grained access control [129778631843866†L57-L58] .
****Simple API****	Developer-friendly API and tooling; multi-tenancy support [129778631843866†L60-L64] .

****Pros****

- * Excellent performance, low latency and good streaming semantics.
- * Durable storage with replay and exactly-once delivery options.
- * Strong security features and multi-tenant isolation.
- * Lighter footprint than Kafka; can run on a small cluster.

****Cons****

- * Operational complexity increases when managing clusters and manual sharding [129778631843866†L159-L165] .
- * Smaller community and ecosystem compared with Kafka; fewer Python client libraries.
- * Still primarily a messaging system; event sourcing requires conventions for stream naming and concurrency, and there is no built-in per-stream versioning.

Apache Pulsar

****Description:**** Pulsar is a multi-tenant, high-performance distributed messaging and streaming platform. It supports both persistent and non-persistent topics and scales by partitioning topics across brokers [541665791063620†L70-L114] .

****Key capabilities****

Capability	Notes
Multi-tenant architecture	Pulsar supports tenants and namespaces with strong isolation [541665791063620†L70-L74] .
Persistent vs non-persistent topics	Messages can be durably stored on disk or kept in memory for transient communication [541665791063620†L96-L110] .
Flexible routing modes	Producers can route messages via round-robin, single partition or custom partitioning [541665791063620†L131-L146] .
Subscription modes	Consumers can use exclusive, failover or shared subscriptions for different delivery guarantees [541665791063620†L148-L187] .

****Pros****

- * High performance and tiered storage; supports very large backlogs.
- * Built-in multi-tenancy and flexible subscription models.

****Cons****

- * Operationally complex: requires running brokers and Apache BookKeeper; heavier

than Postgres or JetStream.

* For event sourcing, retrieving all events for a specific stream still requires scanning partitions; concurrency control must be implemented externally.

Comparison and Recommendations

Suitability for Janus	Durability & Concurrency	Performance	Security & Ops
EventStoreDB	Immutable events, per-stream indexing and optimistic concurrency [640399355648082†L148-L170]	High write throughput (~15 k writes/s) and real-time streaming [640399355648082†L148-L176]	Built-in TLS, authentication and clustering; separate service to operate Purpose-built for event sourcing; ideal for high-volume or multi-node deployments but adds a new database to manage
PostgreSQL event table	ACID transactions ensure consistency; version implemented via transactions [311468078120902†L508-L520]	Adequate for moderate workloads; indexes on `stream_id` support fast reads Mature security features (TLS, row-level privileges); reuses existing infrastructure Simple and low-overhead; best fit if event volume is moderate and we prefer to avoid operating additional services	
Kafka	Durable log; partitions limit per-stream retrieval; no built-in versioning [640399355648082†L190-L204] [640399355648082†L218-L223]	Extremely high throughput; built for large clusters Supports TLS and ACLs; heavy operations requiring ZooKeeper (or KRaft) Overkill for Janus; poor per-stream access and concurrency control	
NATS JetStream	Durable streams with replay; lacks explicit versioning Low latency (<20 ms) and high throughput [129778631843866†L42-L45]	TLS and fine-grained permissions; manual sharding adds complexity [129778631843866†L159-L165]	Good for real-time streaming and microservices; extra service to manage and smaller ecosystem
Apache Pulsar	Durable topics; partitions used for streams; no per-stream versioning High throughput; tiered storage; multi-tenant [541665791063620†L70-L114]	TLS and multi-tenant isolation; complex to run Suitable for large multi-tenant streaming platforms, not necessary for Janus	

Recommendation

Given Janus' **offline-first**, self-hosted ethos and moderate scale, **implementing the event store in the existing PostgreSQL database** is the most pragmatic choice. Postgres already serves as our relational store, and adding an `events` table with `stream_id`, `version` and `data` fields provides:

- * **Low operational overhead** – no additional database to run; uses the same backup/restore and security procedures.
- * **ACID durability** – ensures that events are persisted reliably and that concurrent writers cannot corrupt streams [311468078120902†L508-L520] .
- * **Security** – Postgres offers mature TLS support, role-based access and full-disk encryption, reducing the attack surface.
- * **Sufficient performance** – for a solo developer or small cluster, Postgres can handle thousands of events per second with proper indexing. If event volume becomes a bottleneck, the architecture can evolve toward EventStoreDB or JetStream.

However, if future workloads demand **high write throughput** or **large-scale event streaming**, **EventStoreDB** would be the next contender. It offers per-stream indexes, optimistic concurrency and real-time subscriptions out of the box [640399355648082†L148-L176] . For purely streaming workloads (e.g., IoT sensor feeds), **NATS JetStream** provides low latency and durable message replay [129778631843866†L46-L55] , but it would require additional operational

effort and doesn't natively provide per-stream versioning.

In summary, ****start with a PostgreSQL-based event store**** and revisit specialized solutions if scaling requirements outgrow the relational approach. This path aligns with Janus' modular, offline-first design and minimizes complexity while preserving future flexibility.